

Chapter 1

Language Modeling

(Course notes for NLP by Michael Collins, Columbia University)

1.1 Introduction

In this chapter we will consider the the problem of constructing a *language model* from a set of example sentences in a language. Language models were originally developed for the problem of speech recognition; they still play a central role in modern speech recognition systems. They are also widely used in other NLP applications. The parameter estimation techniques that were originally developed for language modeling, as described in this chapter, are useful in many other contexts, such as the tagging and parsing problems considered in later chapters of this book.

Our task is as follows. Assume that we have a *corpus*, which is a set of sentences in some language. For example, we might have several years of text from the New York Times, or we might have a very large amount of text from the web. Given this corpus, we'd like to estimate the parameters of a language model.

A language model is defined as follows. First, we will define \mathcal{V} to be the set of all words in the language. For example, when building a language model for English we might have

$$\mathcal{V} = \{\text{the, dog, laughs, saw, barks, cat, } \dots\}$$

In practice \mathcal{V} can be quite large: it might contain several thousands, or tens of thousands, of words. We assume that \mathcal{V} is a finite set. A sentence in the language is a sequence of words

$$x_1 x_2 \dots x_n$$

where the integer n is such that $n \geq 1$, we have $x_i \in \mathcal{V}$ for $i \in \{1 \dots (n - 1)\}$, and we assume that x_n is a special symbol, STOP (we assume that STOP is not a

member of \mathcal{V}). We'll soon see why it is convenient to assume that each sentence ends in the STOP symbol. Example sentences could be

the dog barks STOP
 the cat laughs STOP
 the cat saw the dog STOP
 the STOP
 cat the dog the STOP
 cat cat cat STOP
 STOP
 ...

We will define \mathcal{V}^\dagger to be the set of all sentences with the vocabulary \mathcal{V} : this is an infinite set, because sentences can be of any length.

We then give the following definition:

Definition 1 (Language Model) A language model consists of a finite set \mathcal{V} , and a function $p(x_1, x_2, \dots, x_n)$ such that:

1. For any $\langle x_1 \dots x_n \rangle \in \mathcal{V}^\dagger$, $p(x_1, x_2, \dots, x_n) \geq 0$
2. In addition,

$$\sum_{\langle x_1 \dots x_n \rangle \in \mathcal{V}^\dagger} p(x_1, x_2, \dots, x_n) = 1$$

Hence $p(x_1, x_2, \dots, x_n)$ is a probability distribution over the sentences in \mathcal{V}^\dagger .

As one example of a (very bad) method for learning a language model from a training corpus, consider the following. Define $c(x_1 \dots x_n)$ to be the number of times that the sentence $x_1 \dots x_n$ is seen in our training corpus, and N to be the total number of sentences in the training corpus. We could then define

$$p(x_1 \dots x_n) = \frac{c(x_1 \dots x_n)}{N}$$

This is, however, a very poor model: in particular it will assign probability 0 to any sentence not seen in the training corpus. Thus it fails to generalize to sentences that have not been seen in the training data. The key technical contribution of this chapter will be to introduce methods that do generalize to sentences that are not seen in our training data.

At first glance the language modeling problem seems like a rather strange task, so why are we considering it? There are a couple of reasons:

1. Language models are very useful in a broad range of applications, the most obvious perhaps being speech recognition and machine translation. In many applications it is very useful to have a good “prior” distribution $p(x_1 \dots x_n)$ over which sentences are or aren’t probable in a language. For example, in speech recognition the language model is combined with an acoustic model that models the pronunciation of different words: one way to think about it is that the acoustic model generates a large number of candidate sentences, together with probabilities; the language model is then used to reorder these possibilities based on how likely they are to be a sentence in the language.
2. The techniques we describe for defining the function p , and for estimating the parameters of the resulting model from training examples, will be useful in several other contexts during the course: for example in hidden Markov models, which we will see next, and in models for natural language parsing.

1.2 Markov Models

We now turn to a critical question: given a training corpus, how do we learn the function p ? In this section we describe *Markov models*, a central idea from probability theory; in the next section we describe *trigram language models*, an important class of language models that build directly on ideas from Markov models.

1.2.1 Markov Models for Fixed-length Sequences

Consider a sequence of random variables, X_1, X_2, \dots, X_n . Each random variable can take any value in a finite set \mathcal{V} . For now we will assume that the length of the sequence, n , is some fixed number (e.g., $n = 100$). In the next section we’ll describe how to generalize the approach to cases where n is also a random variable, allowing different sequences to have different lengths.

Our goal is as follows: we would like to model the probability of any sequence $x_1 \dots x_n$, where $n \geq 1$ and $x_i \in \mathcal{V}$ for $i = 1 \dots n$, that is, to model the joint probability

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n)$$

There are $|\mathcal{V}|^n$ possible sequences of the form $x_1 \dots x_n$: so clearly, it is not feasible for reasonable values of $|\mathcal{V}|$ and n to simply list all $|\mathcal{V}|^n$ probabilities. We would like to build a much more compact model.

In a **first-order Markov process**, we make the following assumption, which considerably simplifies the model:

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n)$$

$$= P(X_1 = x_1) \prod_{i=2}^n P(X_i = x_i | X_1 = x_1, \dots, X_{i-1} = x_{i-1}) \quad (1.1)$$

$$= P(X_1 = x_1) \prod_{i=2}^n P(X_i = x_i | X_{i-1} = x_{i-1}) \quad (1.2)$$

The first step, in Eq. 1.1, is exact: by the chain rule of probabilities, *any* distribution $P(X_1 = x_1 \dots X_n = x_n)$ can be written in this form. So we have made no assumptions in this step of the derivation. However, the second step, in Eq. 1.2, is not necessarily exact: we have made the assumption that for any $i \in \{2 \dots n\}$, for any $x_1 \dots x_i$,

$$P(X_i = x_i | X_1 = x_1 \dots X_{i-1} = x_{i-1}) = P(X_i = x_i | X_{i-1} = x_{i-1})$$

This is a (first-order) *Markov assumption*. We have assumed that the identity of the i 'th word in the sequence depends only on the identity of the previous word, x_{i-1} . More formally, we have assumed that the value of X_i is conditionally independent of $X_1 \dots X_{i-2}$, given the value for X_{i-1} .

In a *second-order Markov process*, which will form the basis of trigram language models, we make a slightly weaker assumption, namely that each word depends on the previous *two* words in the sequence:

$$\begin{aligned} & P(X_i = x_i | X_1 = x_1, \dots, X_{i-1} = x_{i-1}) \\ &= P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1}) \end{aligned}$$

It follows that the probability of an entire sequence is written as

$$\begin{aligned} & P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) \\ &= \prod_{i=1}^n P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1}) \end{aligned} \quad (1.3)$$

For convenience, we will assume that $x_0 = x_{-1} = *$ in this definition, where $*$ is a special “start” symbol in the sentence.

1.2.2 Markov Sequences for Variable-length Sentences

In the previous section, we assumed that the length of the sequence, n , was fixed. In many applications, however, the length n can itself vary. Thus n is itself a random variable. There are various ways of modeling this variability in length: in this section we describe the most common approach for language modeling.

The approach is simple: we will assume that the n 'th word in the sequence, X_n , is always equal to a special symbol, the STOP symbol. This symbol can only

appear at the end of a sequence. We use exactly the same assumptions as before: for example under a second-order Markov assumption, we have

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) = \prod_{i=1}^n P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1}) \quad (1.4)$$

for any $n \geq 1$, and for any $x_1 \dots x_n$ such that $x_n = \text{STOP}$, and $x_i \in \mathcal{V}$ for $i = 1 \dots (n-1)$.

We have assumed a second-order Markov process where at each step we generate a symbol x_i from the distribution

$$P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1})$$

where x_i can be a member of \mathcal{V} , or alternatively can be the STOP symbol. If we generate the STOP symbol, we finish the sequence. Otherwise, we generate the next symbol in the sequence.

A little more formally, the process that generates sentences would be as follows:

1. Initialize $i = 1$, and $x_0 = x_{-1} = *$
2. Generate x_i from the distribution

$$P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1})$$

3. If $x_i = \text{STOP}$ then return the sequence $x_1 \dots x_i$. Otherwise, set $i = i + 1$ and return to step 2.

Thus we now have a model that generates sequences that vary in length.

1.3 Trigram Language Models

There are various ways of defining language models, but we'll focus on a particularly important example, the trigram language model, in this chapter. This will be a direct application of Markov models, as described in the previous section, to the language modeling problem. In this section we give the basic definition of a trigram model, discuss maximum-likelihood parameter estimates for trigram models, and finally discuss strengths of weaknesses of trigram models.

1.3.1 Basic Definitions

As in Markov models, we model each sentence as a sequence of n random variables, X_1, X_2, \dots, X_n . The length, n , is itself a random variable (it can vary across different sentences). We always have $X_n = \text{STOP}$. Under a second-order Markov model, the probability of any sentence $x_1 \dots x_n$ is then

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) = \prod_{i=1}^n P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1})$$

where we assume as before that $x_0 = x_{-1} = *$.

We will assume that for any i , for any x_{i-2}, x_{i-1}, x_i ,

$$P(X_i = x_i | X_{i-2} = x_{i-2}, X_{i-1} = x_{i-1}) = q(x_i | x_{i-2}, x_{i-1})$$

where $q(w|u, v)$ for any (u, v, w) is a parameter of the model. We will soon see how to derive estimates of the $q(w|u, v)$ parameters from our training corpus. Our model then takes the form

$$p(x_1 \dots x_n) = \prod_{i=1}^n q(x_i | x_{i-2}, x_{i-1})$$

for any sequence $x_1 \dots x_n$.

This leads us to the following definition:

Definition 2 (Trigram Language Model) *A trigram language model consists of a finite set \mathcal{V} , and a parameter*

$$q(w|u, v)$$

for each trigram u, v, w such that $w \in \mathcal{V} \cup \{\text{STOP}\}$, and $u, v \in \mathcal{V} \cup \{\}$. The value for $q(w|u, v)$ can be interpreted as the probability of seeing the word w immediately after the bigram (u, v) . For any sentence $x_1 \dots x_n$ where $x_i \in \mathcal{V}$ for $i = 1 \dots (n-1)$, and $x_n = \text{STOP}$, the probability of the sentence under the trigram language model is*

$$p(x_1 \dots x_n) = \prod_{i=1}^n q(x_i | x_{i-2}, x_{i-1})$$

where we define $x_0 = x_{-1} = *$. \square

For example, for the sentence

the dog barks STOP

we would have

$$p(\text{the dog barks STOP}) = q(\text{the}|\ast, \ast) \times q(\text{dog}|\ast, \text{the}) \times q(\text{barks}|\text{the, dog}) \times q(\text{STOP}|\text{dog, barks})$$

Note that in this expression we have one term for each word in the sentence (*the*, *dog*, *barks*, and *STOP*). Each word depends only on the previous two words: this is the trigram assumption.

The parameters satisfy the constraints that for any trigram u, v, w ,

$$q(w|u, v) \geq 0$$

and for any bigram u, v ,

$$\sum_{w \in \mathcal{V} \cup \{\text{STOP}\}} q(w|u, v) = 1$$

Thus $q(w|u, v)$ defines a distribution over possible words w , conditioned on the bigram context u, v .

The key problem we are left with is to estimate the parameters of the model, namely

$$q(w|u, v)$$

where w can be any member of $\mathcal{V} \cup \{\text{STOP}\}$, and $u, v \in \mathcal{V} \cup \{\ast\}$. There are around $|\mathcal{V}|^3$ parameters in the model. This is likely to be a very large number. For example with $|\mathcal{V}| = 10,000$ (this is a realistic number, most likely quite small by modern standards), we have $|\mathcal{V}|^3 \approx 10^{12}$.

1.3.2 Maximum-Likelihood Estimates

We first start with the most generic solution to the estimation problem, the *maximum-likelihood estimates*. We will see that these estimates are flawed in a critical way, but we will then show how related estimates can be derived that work very well in practice.

First, some notation. Define $c(u, v, w)$ to be the number of times that the trigram (u, v, w) is seen in the training corpus: for example, $c(\text{the, dog, barks})$ is the number of times that the sequence of three words *the dog barks* is seen in the training corpus. Similarly, define $c(u, v)$ to be the number of times that the bigram (u, v) is seen in the corpus. For any w, u, v , we then define

$$q(w|u, v) = \frac{c(u, v, w)}{c(u, v)}$$

As an example, our estimate for $q(\text{barks}|\text{the, dog})$ would be

$$q(\text{barks}|\text{the, dog}) = \frac{c(\text{the, dog, barks})}{c(\text{the, dog})}$$

This estimate is very natural: the numerator is the number of times the entire trigram *the dog barks* is seen, and the denominator is the number of times the bigram *the dog* is seen. We simply take the ratio of these two terms.

Unfortunately, this way of estimating parameters runs into a very serious issue. Recall that we have a very large number of parameters in our model (e.g., with a vocabulary size of 10,000, we have around 10^{12} parameters). Because of this, many of our counts will be zero. This leads to two problems:

- Many of the above estimates will be $q(w|u, v) = 0$, due to the count in the numerator being 0. This will lead to many trigram probabilities being systematically underestimated: it seems unreasonable to assign probability 0 to any trigram not seen in training data, given that the number of parameters of the model is typically very large in comparison to the number of words in the training corpus.
- In cases where the denominator $c(u, v)$ is equal to zero, the estimate is not well defined.

We will shortly see how to come up with modified estimates that fix these problems. First, however, we discuss how language models are evaluated, and then discuss strengths and weaknesses of trigram language models.

1.3.3 Evaluating Language Models: Perplexity

So how do we measure the quality of a language model? A very common method is to evaluate the *perplexity* of the model on some held-out data.

The method is as follows. Assume that we have some test data sentences $x^{(1)}, x^{(2)}, \dots, x^{(m)}$. Each test sentence $x^{(i)}$ for $i \in \{1 \dots m\}$ is a sequence of words $x_1^{(i)}, \dots, x_{n_i}^{(i)}$, where n_i is the length of the i 'th sentence. As before we assume that every sentence ends in the STOP symbol.

It is critical that the test sentences are “held out”, in the sense that *they are not part of the corpus used to estimate the language model*. In this sense, they are examples of new, unseen sentences.

For any test sentence $x^{(i)}$, we can measure its probability $p(x^{(i)})$ under the language model. A natural measure of the quality of the language model would be

the probability it assigns to the entire set of test sentences, that is

$$\prod_{i=1}^m p(x^{(i)})$$

The intuition is as follows: the higher this quantity is, the better the language model is at modeling unseen sentences.

The perplexity on the test corpus is derived as a direct transformation of this quantity. Define M to be the total number of words in the test corpus. More precisely, under the definition that n_i is the length of the i 'th test sentence,

$$M = \sum_{i=1}^m n_i$$

Then the average log probability under the model is defined as

$$\frac{1}{M} \log_2 \prod_{i=1}^m p(x^{(i)}) = \frac{1}{M} \sum_{i=1}^m \log_2 p(x^{(i)})$$

This is just the log probability of the entire test corpus, divided by the total number of words in the test corpus. Here we use $\log_2(z)$ for any $z > 0$ to refer to the log with respect to base 2 of z . Again, the higher this quantity is, the better the language model.

The *perplexity* is then defined as

$$2^{-l}$$

where

$$l = \frac{1}{M} \sum_{i=1}^m \log_2 p(x^{(i)})$$

Thus we take the negative of the average log probability, and raise two to that power. (Again, we're assuming in this section that \log_2 is log base two). The perplexity is a positive number. The *smaller* the value of perplexity, the better the language model is at modeling unseen data.

Some intuition behind perplexity is as follows. Say we have a vocabulary \mathcal{V} , where $|\mathcal{V} \cup \{\text{STOP}\}| = N$, and the model predicts

$$q(w|u, v) = \frac{1}{N}$$

for all u, v, w . Thus this is the dumb model that simply predicts the uniform distribution over the vocabulary together with the STOP symbol. In this case, it can

be shown that the perplexity is equal to N . So under a uniform probability model, the perplexity is equal to the vocabulary size. Perplexity can be thought of as the effective vocabulary size under the model: if, for example, the perplexity of the model is 120 (even though the vocabulary size is say 10,000), then this is roughly equivalent to having an effective vocabulary size of 120.

To give some more motivation, it is relatively easy to show that the perplexity is equal to

$$\frac{1}{t}$$

where

$$t = \sqrt[M]{\prod_{i=1}^m p(x^{(i)})}$$

Here we use $\sqrt[M]{z}$ to refer to the M 'th root of z : so t is the value such that $t^M = \prod_{i=1}^m p(x^{(i)})$. Given that

$$\prod_{i=1}^m p(x^{(i)}) = \prod_{i=1}^m \prod_{j=1}^{n_i} q(x_j^{(i)} | x_{j-2}^{(i)}, x_{j-1}^{(i)})$$

and $M = \sum_i n_i$, the value for t is the *geometric mean* of the terms $q(x_j^{(i)} | x_{j-2}^{(i)}, x_{j-1}^{(i)})$ appearing in $\prod_{i=1}^m p(x^{(i)})$. For example if the perplexity is equal to 100, then $t = 0.01$, indicating that the geometric mean is 0.01.

One additional useful fact about perplexity is the following. If for any trigram u, v, w seen in test data, we have the estimate

$$q(w|u, v) = 0$$

then the perplexity will be ∞ . To see this, note that in this case the probability of the test corpus under the model will be 0, and the average log probability will be $-\infty$. Thus if we take perplexity seriously as our measure of a language model, then we should avoid giving 0 estimates at all costs.

Finally, some intuition about “typical” values for perplexity. Goodman (“A bit of progress in language modeling”, figure 2) evaluates unigram, bigram and trigram language models on English data, with a vocabulary size of 50,000. In a bigram model we have parameters of the form $q(w|v)$, and

$$p(x_1 \dots x_n) = \prod_{i=1}^n q(x_i | x_{i-1})$$

Thus each word depends only on the previous word in the sentence. In a unigram model we have parameters $q(w)$, and

$$p(x_1 \dots x_n) = \prod_{i=1}^n q(x_i)$$

Thus each word is chosen completely independently of other words in the sentence. Goodman reports perplexity figures of around 74 for a trigram model, 137 for a bigram model, and 955 for a unigram model. The perplexity for a model that simply assigns probability $1/50,000$ to each word in the vocabulary would be 50,000. So the trigram model clearly gives a big improvement over bigram and unigram models, and a huge improvement over assigning a probability of $1/50,000$ to each word in the vocabulary.

1.3.4 Strengths and Weaknesses of Trigram Language Models

The trigram assumption is arguably quite strong, and linguistically naive (see the lecture slides for discussion). However, it leads to models that are very useful in practice.

1.4 Smoothed Estimation of Trigram Models

As discussed previously, a trigram language model has a very large number of parameters. The maximum-likelihood parameter estimates, which take the form

$$q(w|u, v) = \frac{c(u, v, w)}{c(u, v)}$$

will run into serious issues with sparse data. Even with a large set of training sentences, many of the counts $c(u, v, w)$ or $c(u, v)$ will be low, or will be equal to zero.

In this section we describe *smoothed estimation methods*, which alleviate many of the problems found with sparse data. The key idea will be to rely on lower-order statistical estimates—in particular, estimates based on bigram or unigram counts—to “smooth” the estimates based on trigrams. We discuss two smoothing methods that are very commonly used in practice: first, *linear interpolation*; second, *discounting methods*.

1.4.1 Linear Interpolation

A linearly interpolated trigram model is derived as follows. We define the *trigram*, *bigram*, and *unigram* maximum-likelihood estimates as

$$q_{ML}(w|u, v) = \frac{c(u, v, w)}{c(u, v)}$$

$$q_{ML}(w|v) = \frac{c(v, w)}{c(v)}$$

$$q_{ML}(w) = \frac{c(w)}{c()}$$

where we have extended our notation: $c(w)$ is the number of times word w is seen in the training corpus, and $c()$ is the total number of words seen in the training corpus.

The trigram, bigram, and unigram estimates have different strengths and weaknesses. The unigram estimate will never have the problem of its numerator or denominator being equal to 0: thus the estimate will always be well-defined, and will always be greater than 0 (providing that each word is seen at least once in the training corpus, which is a reasonable assumption). However, the unigram estimate completely ignores the context (previous two words), and hence discards very valuable information. In contrast, the trigram estimate does make use of context, but has the problem of many of its counts being 0. The bigram estimate falls between these two extremes.

The idea in linear interpolation is to use all three estimates, by defining the trigram estimate as follows:

$$q(w|u, v) = \lambda_1 \times q_{ML}(w|u, v) + \lambda_2 \times q_{ML}(w|v) + \lambda_3 \times q_{ML}(w)$$

Here λ_1 , λ_2 and λ_3 are three additional parameters of the model, which satisfy

$$\lambda_1 \geq 0, \lambda_2 \geq 0, \lambda_3 \geq 0$$

and

$$\lambda_1 + \lambda_2 + \lambda_3 = 1$$

Thus we take a weighted average of the three estimates.

There are various ways of estimating the λ values. A common one is as follows. Say we have some additional held-out data, which is separate from both our training and test corpora. We will call this data the *development data*. Define

$c'(u, v, w)$ to be the number of times that the trigram (u, v, w) is seen in the development data. It is easy to show that the log-likelihood of the development data, as a function of the parameters $\lambda_1, \lambda_2, \lambda_3$, is

$$\begin{aligned} L(\lambda_1, \lambda_2, \lambda_3) &= \sum_{u,v,w} c'(u, v, w) \log q(w|u, v) \\ &= \sum_{u,v,w} c'(u, v, w) \log (\lambda_1 \times q_{ML}(w|u, v) + \lambda_2 \times q_{ML}(w|v) + \lambda_3 \times q_{ML}(w)) \end{aligned}$$

We would like to choose our λ values to make $L(\lambda_1, \lambda_2, \lambda_3)$ as high as possible. Thus the λ values are taken to be

$$\arg \max_{\lambda_1, \lambda_2, \lambda_3} L(\lambda_1, \lambda_2, \lambda_3)$$

such that

$$\lambda_1 \geq 0, \lambda_2 \geq 0, \lambda_3 \geq 0$$

and

$$\lambda_1 + \lambda_2 + \lambda_3 = 1$$

Finding the optimal values for $\lambda_1, \lambda_2, \lambda_3$ is fairly straightforward (see section ?? for one algorithm that is often used for this purpose).

As described, our method has three smoothing parameters, λ_1, λ_2 , and λ_3 . The three parameters can be interpreted as an indication of the confidence, or weight, placed on each of the trigram, bigram, and unigram estimates. For example, if λ_1 is close to 1, this implies that we put a significant weight on the trigram estimate $q_{ML}(w|u, v)$; conversely, if λ_1 is close to zero we have placed a low weighting on the trigram estimate.

In practice, it is important to add an additional degree of freedom, by allowing the values for λ_1, λ_2 and λ_3 to vary depending on the bigram (u, v) that is being conditioned on. In particular, the method can be extended to allow λ_1 to be larger when $c(u, v)$ is larger—the intuition being that a larger value of $c(u, v)$ should translate to us having more confidence in the trigram estimate.

At the very least, the method is used to ensure that $\lambda_1 = 0$ when $c(u, v) = 0$, because in this case the trigram estimate

$$q_{ML}(w|u, v) = \frac{c(u, v, w)}{c(u, v)}$$

is undefined. Similarly, if both $c(u, v)$ and $c(v)$ are equal to zero, we need $\lambda_1 = \lambda_2 = 0$, as both the trigram and bigram ML estimates are undefined.

One extension to the method, often referred to as *bucketing*, is described in section 1.5.1. Another, much simpler method, is to define

$$\begin{aligned}\lambda_1 &= \frac{c(u, v)}{c(u, v) + \gamma} \\ \lambda_2 &= (1 - \lambda_1) \times \frac{c(v)}{c(v) + \gamma} \\ \lambda_3 &= 1 - \lambda_1 - \lambda_2\end{aligned}$$

where $\gamma > 0$ is the only parameter of the method. It can be verified that $\lambda_1 \geq 0$, $\lambda_2 \geq 0$, and $\lambda_3 \geq 0$, and also that $\lambda_1 + \lambda_2 + \lambda_3 = 1$.

Under this definition, it can be seen that λ_1 increases as $c(u, v)$ increases, and similarly that λ_2 increases as $c(v)$ increases. In addition we have $\lambda_1 = 0$ if $c(u, v) = 0$, and $\lambda_2 = 0$ if $c(v) = 0$. The value for γ can again be chosen by maximizing log-likelihood of a set of development data.

This method is relatively crude, and is not likely to be optimal. It is, however, very simple, and in practice it can work well in some applications.

1.4.2 Discounting Methods

We now describe an alternative estimation method, which is again commonly used in practice. Consider first a method for estimating a bigram language model, that is, our goal is to define

$$q(w|v)$$

for any $w \in \mathcal{V} \cup \{\text{STOP}\}$, $v \in \mathcal{V} \cup \{*\}$.

The first step will be to define *discounted counts*. For any bigram $c(v, w)$ such that $c(v, w) > 0$, we define the discounted count as

$$c^*(v, w) = c(v, w) - \beta$$

where β is a value between 0 and 1 (a typical value might be $\beta = 0.5$). Thus we simply subtract a constant value, β , from the count. This reflects the intuition that if we take counts from the training corpus, we will systematically over-estimate the probability of bigrams seen in the corpus (and under-estimate bigrams not seen in the corpus).

For any bigram (v, w) such that $c(v, w) > 0$, we can then define

$$q(w|v) = \frac{c^*(v, w)}{c(v)}$$

Thus we use the discounted count on the numerator, and the regular count on the denominator of this expression.

x	$c(x)$	$c^*(x)$	$\frac{c^*(x)}{c(the)}$
the	48		
the, dog	15	14.5	14.5/48
the, woman	11	10.5	10.5/48
the, man	10	9.5	9.5/48
the, park	5	4.5	4.5/48
the, job	2	1.5	1.5/48
the, telescope	1	0.5	0.5/48
the, manual	1	0.5	0.5/48
the, afternoon	1	0.5	0.5/48
the, country	1	0.5	0.5/48
the, street	1	0.5	0.5/48

Figure 1.1: An example illustrating discounting methods. We show a made-up example, where the unigram *the* is seen 48 times, and we list all bigrams (u, v) such that $u = the$, and $c(u, v) > 0$ (the counts for these bigrams sum to 48). In addition we show the discounted count $c^*(x) = c(x) - \beta$, where $\beta = 0.5$, and finally we show the estimate $c^*(x)/c(the)$ based on the discounted count.

For any context v , this definition leads to some *missing probability mass*, defined as

$$\alpha(v) = 1 - \sum_{w:c(v,w)>0} \frac{c^*(v,w)}{c(v)}$$

As an example, consider the counts shown in the example in figure 1.1. In this case we show all bigrams (u, v) where $u = \text{the}$, and $c(u, v) > 0$. We use a discounting value of $\beta = 0.5$. In this case we have

$$\sum_{w:c(v,w)>0} \frac{c^*(v,w)}{c(v)} = \frac{14.5}{48} + \frac{10.5}{48} + \frac{9.5}{48} + \frac{4.5}{48} + \frac{1.5}{48} + 5 \times \frac{0.5}{48} = \frac{43}{48}$$

and the missing mass is

$$\alpha(\text{the}) = 1 - \frac{43}{48} = \frac{5}{48}$$

The intuition behind discounted methods is to divide this “missing mass” between the words w such that $c(v, w) = 0$.

More specifically, the complete definition of the estimate is as follows. For any v , define the sets

$$\mathcal{A}(v) = \{w : c(v, w) > 0\}$$

and

$$\mathcal{B}(v) = \{w : c(v, w) = 0\}$$

For the data in figure 1.1, for example, we would have

$$\mathcal{A}(\text{the}) = \{\text{dog}, \text{woman}, \text{man}, \text{park}, \text{job}, \text{telescope}, \text{manual}, \text{afternoon}, \text{country}, \text{street}\}$$

and $\mathcal{B}(\text{the})$ would be the set of remaining words in the vocabulary.

Then the estimate is defined as

$$q_D(w|v) = \begin{cases} \frac{c^*(v,w)}{c(v)} & \text{If } w \in \mathcal{A}(v) \\ \alpha(v) \times \frac{q_{ML}(w)}{\sum_{w \in \mathcal{B}(v)} q_{ML}(w)} & \text{If } w \in \mathcal{B}(v) \end{cases}$$

Thus if $c(v, w) > 0$ we return the estimate $c^*(v, w)/c(v)$; otherwise we divide the remaining probability mass $\alpha(v)$ in proportion to the unigram estimates $q_{ML}(w)$.

The method can be generalized to trigram language models in a natural, recursive way: for any bigram (u, v) define

$$\mathcal{A}(u, v) = \{w : c(u, v, w) > 0\}$$

and

$$\mathcal{B}(u, v) = \{w : c(u, v, w) = 0\}$$

Define $c^*(u, v, w)$ to be the discounted count for the trigram (u, v, w) : that is,

$$c^*(u, v, w) = c(u, v, w) - \beta$$

where β is again the discounting value. Then the trigram model is

$$q_D(w|u, v) = \begin{cases} \frac{c^*(u, v, w)}{c(u, v)} & \text{If } w \in \mathcal{A}(u, v) \\ \alpha(u, v) \times \frac{q_D(w|v)}{\sum_{w \in \mathcal{B}(u, v)} q_D(w|v)} & \text{If } w \in \mathcal{B}(u, v) \end{cases}$$

where

$$\alpha(u, v) = 1 - \sum_{w \in \mathcal{A}(u, v)} \frac{c^*(u, v, w)}{c(u, v)}$$

is again the “missing” probability mass. Note that we have divided the missing probability mass in proportion to the bigram estimates $q_D(w|v)$, which were defined previously.

The only parameter of this approach is the discounting value, β . As in linearly interpolated models, the usual way to choose the value for this parameter is by optimization of likelihood of a development corpus, which is again a separate set of data from the training and test corpora. Define $c'(u, v, w)$ to be the number of times that the trigram u, v, w is seen in this development corpus. The log-likelihood of the development data is

$$\sum_{u, v, w} c'(u, v, w) \log q_D(w|u, v)$$

where $q_D(w|u, v)$ is defined as above. The parameter estimates $q_D(w|u, v)$ will vary as the value for β varies. Typically we will test a set of possible values for β —for example, we might test all values in the set $\{0.1, 0.2, 0.3, \dots, 0.9\}$ —where for each value of β we calculate the log-likelihood of the development data. We then choose the value for β that maximizes this log-likelihood.

1.5 Advanced Topics

1.5.1 Linear Interpolation with Bucketing

In linearly interpolated models the parameter estimates are defined as

$$q(w|u, v) = q(w|u, v) = \lambda_1 \times q_{ML}(w|u, v) + \lambda_2 \times q_{ML}(w|v) + \lambda_3 \times q_{ML}(w)$$

where λ_1, λ_2 and λ_3 are smoothing parameters in the approach.

In practice, it is important to allow the smoothing parameters to vary depending on the bigram (u, v) that is being conditioned on—in particular, the higher the count $c(u, v)$, the higher the weight should be for λ_1 (and similarly, the higher the count $c(v)$, the higher the weight should be for λ_2). The classical method for achieving this is through an extension that is often referred to as “bucketing”.

The first step in this method is to define a function Π that maps bigrams (u, v) to values $\Pi(u, v) \in \{1, 2, \dots, K\}$ where K is some integer specifying the number of buckets. Thus the function Π defines a partition of bigrams into K different subsets. The function is defined by hand, and typically depends on the counts seen in the training data. One such definition, with $K = 3$, would be

$$\begin{aligned}\Pi(u, v) &= 1 && \text{if } c(u, v) > 0 \\ \Pi(u, v) &= 2 && \text{if } c(u, v) = 0 \text{ and } c(v) > 0 \\ \Pi(u, v) &= 3 && \text{otherwise}\end{aligned}$$

This is a very simple definition that simply tests whether the counts $c(u, v)$ and $c(v)$ are equal to 0.

Another slightly more complicated definition, which is more sensitive to frequency of the bigram (u, v) , would be

$$\begin{aligned}\Pi(u, v) &= 1 && \text{if } 100 \leq c(u, v) \\ \Pi(u, v) &= 2 && \text{if } 50 \leq c(u, v) < 100 \\ \Pi(u, v) &= 3 && \text{if } 20 \leq c(u, v) < 50 \\ \Pi(u, v) &= 4 && \text{if } 10 \leq c(u, v) < 20 \\ \Pi(u, v) &= 5 && \text{if } 5 \leq c(u, v) < 10 \\ \Pi(u, v) &= 6 && \text{if } 2 \leq c(u, v) < 5 \\ \Pi(u, v) &= 7 && \text{if } c(u, v) = 1 \\ \Pi(u, v) &= 8 && \text{if } c(u, v) = 0 \text{ and } c(v) > 0 \\ \Pi(u, v) &= 9 && \text{otherwise}\end{aligned}$$

Given a definition of the function $\Pi(u, v)$, we then introduce smoothing parameters $\lambda_1^{(k)}, \lambda_2^{(k)}, \lambda_3^{(k)}$ for all $k \in \{1 \dots K\}$. Thus each bucket has its own set of smoothing parameters. We have the constraints that for all $k \in \{1 \dots K\}$

$$\lambda_1^{(k)} \geq 0, \lambda_2^{(k)} \geq 0, \lambda_3^{(k)} \geq 0$$

and

$$\lambda_1^{(k)} + \lambda_2^{(k)} + \lambda_3^{(k)} = 1$$

The linearly interpolated estimate will be

$$q(w|u, v) = q(w|u, v) = \lambda_1^{(k)} \times q_{ML}(w|u, v) + \lambda_2^{(k)} \times q_{ML}(w|v) + \lambda_3^{(k)} \times q_{ML}(w)$$

where

$$k = \Pi(u, v)$$

Thus we have crucially introduced a dependence of the smoothing parameters on the value for $\Pi(u, v)$. Thus each bucket of bigrams gets its own set of smoothing parameters; the values for the smoothing parameters can vary depending on the value of $\Pi(u, v)$ (which is usually directly related to the counts $c(u, v)$ and $c(v)$).

The smoothing parameters are again estimated using a development data set. If we again define $c'(u, v, w)$ to be the number of times the trigram u, v, w appears in the development data, the log-likelihood of the development data is

$$\begin{aligned} & \sum_{u,v,w} c'(u, v, w) \log q(w|u, v) \\ = & \sum_{u,v,w} c'(u, v, w) \log \left(\lambda_1^{(\Pi(u,v))} \times q_{ML}(w|u, v) + \lambda_2^{(\Pi(u,v))} \times q_{ML}(w|v) + \lambda_3^{(\Pi(u,v))} \times q_{ML}(w) \right) \\ = & \sum_{k=1}^K \sum_{u,v,w: \Pi(u,v)=k} c'(u, v, w) \log \left(\lambda_1^{(k)} \times q_{ML}(w|u, v) + \lambda_2^{(k)} \times q_{ML}(w|v) + \lambda_3^{(k)} \times q_{ML}(w) \right) \end{aligned}$$

The $\lambda_1^{(k)}, \lambda_2^{(k)}, \lambda_3^{(k)}$ values are chosen to maximize this function.